

# About Dataset

Customer churn refers to the phenomenon where customers discontinue their relationship or subscription with a company or service provider. It represents the rate at which customers stop using a company's products or services within a specific period. Churn is an important metric for businesses as it directly impacts revenue, growth, and customer retention.

## The dataset contains

- age
- gender
- tenure
- usage frequency
- support calls
- payment delay
- subscription type
- contract length
- total spends
- last interaction

Link – <https://www.kaggle.com/datasets/muhammadshahidazeem/customer-churn-dataset>

## Importing data

Code:

```
data <- read.csv("D:/AIUB/CSE SEM 10/DATA WAREHOUSING AND DATA MINING/final project/data/training.csv")
head(data)
```

output:

	CustomerID	Age	Gender	Tenure	Usage.Frequency	Support.Calls	Payment.Delay	Subscription.Type
1	2	30	Female	39	14	5	18	Standard
2	3	65	Female	49	1	10	8	Basic
3	4	55	Female	14	4	6	18	Basic
4	5	58	Male	38	21	7	7	Standard
5	6	23	Male	32	20	5	8	Basic
6	8	51	Male	33	25	9	26	Premium
	Contract.Length	Total.Spend	Last.Interaction	Churn				
1	Annual	932	17	1				
2	Monthly	557	6	1				
3	Quarterly	185	3	1				
4	Monthly	396	29	1				
5	Monthly	617	20	1				
6	Annual	129	8	1				



Checking for duplicate or missing values in the dataset

**Code:**

```
sum(duplicated(data))
```

```
sum(is.na(data))
```

**Output:**

```
> sum(duplicated(data))  
[1] 0  
> sum(is.na(data))  
[1] 8
```

8 missing values found. So, we decided to delete the rows containing missing values

**Code:**

```
data <- na.omit(data)
```

```
|  
sum(is.na(data))
```

**Output:**


```
> sum(is.na(data))  
[1] 0
```

Showing correlation between total money spend, payment delay and Tenure column

**Code:**

```
subset(as.matrix(Filter(is.numeric, na.exclude(distinct(data)))),  
       select = c(Total.Spend, Payment.Delay, Tenure)) %>% cor() %>% heatmap())
```

**Output:**



### Train Test split

#### Code:

```
sample <- sample(c(TRUE, FALSE), nrow(data), replace=TRUE, prob=c(0.7,0.3))
train <- data[sample, ]
test <- data[!sample, ]
```

## Decision Tree with Gini index

Training the model and Testing

#### Code:

```
DT_Model_gini <- rpart(formula = Churn ~., data = train, method = "class", parms = list(split = "gini"))

test_pred_gini = predict(DT_Model_gini, newdata= test, type = "class")
```

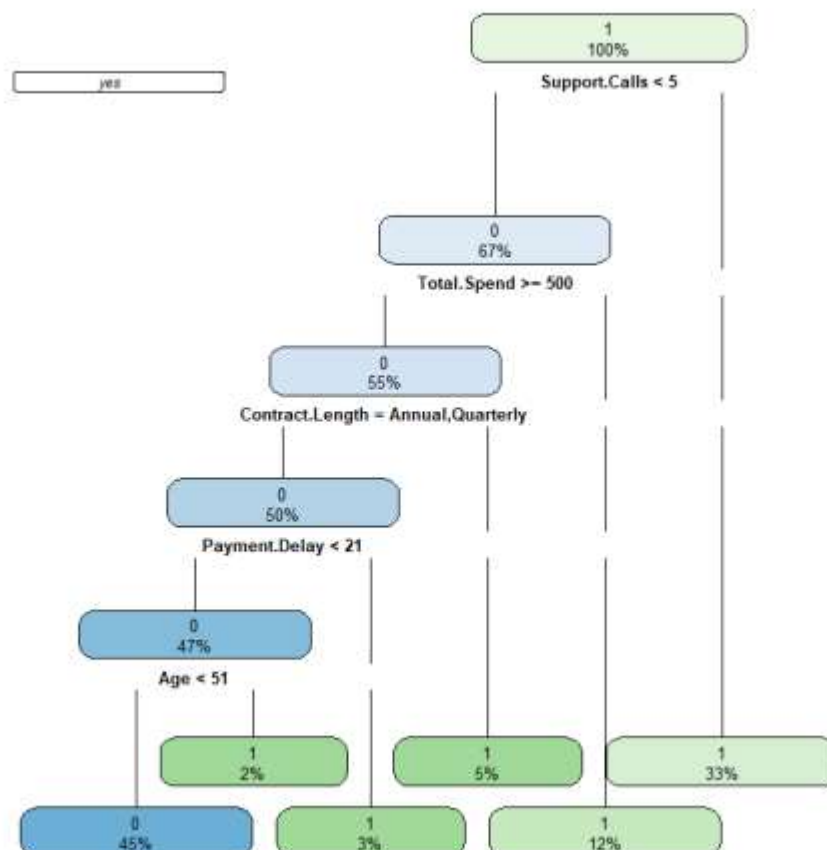
```
[1] 0.9756737
```

```

rpart.plot(rpart(formula = Churn ~., data = train,
                  method = "class", parms = list(split = "gini")), extra = 100)

```

**Output:**



Showing the confusion matrix

**Code:**

```
confusion_matrix <- table(test_pred_gini, test$Churn)
```

```
print(confusion_matrix)
```

**Output:**

```
test_pred_gini      0      1
      0 56701 2813
      1   395 71965
```

Showing Precision, recall and F1 score

**Code:**

```
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
precision <- diag(confusion_matrix) / rowSums(confusion_matrix)
recall <- diag(confusion_matrix) / colSums(confusion_matrix)
f1_score <- 2 * (precision * recall) / (precision + recall)
report <- data.frame(Class = rownames(confusion_matrix),
                     Precision = precision,
                     Recall = recall,
                     F1_Score = f1_score)

print(report)
```

**Output:**

```
Class Precision Recall F1_Score
0      0 0.9527338 0.9930818 0.9724895
1      1 0.9945412 0.9623820 0.9781973
```

## Decision Tree with Information gain

Training the model and Testing

**Code:**

```
DT_Model_info <- rpart(formula = Churn ~., data = train, method = "class", parms = list(split = "information"))  
test_pred_info = predict(DT_Model_info, newdata= test, type = "class")
```

Showing accuracy

**Code:**

```
sum(test_pred_info == test$Churn) / nrow(test)
```

**Output:**

```
0.9871923
```

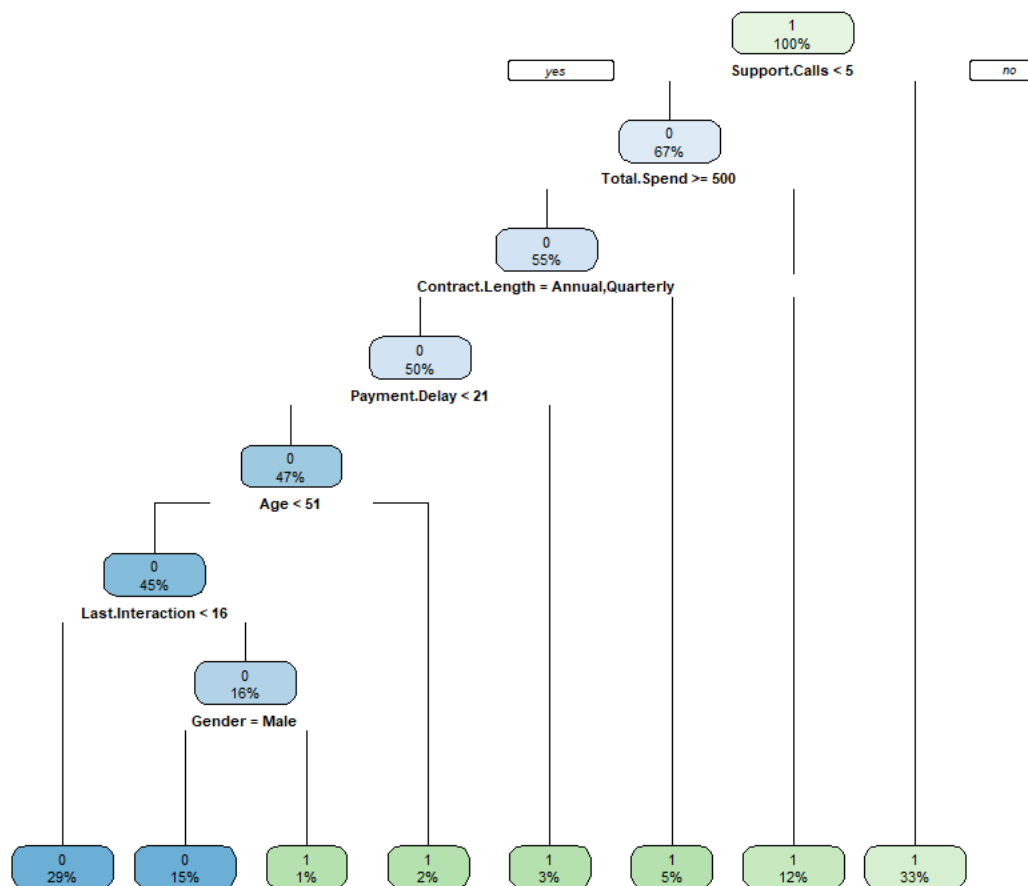
**Code:**

Showing How many rules are generated

**Code:**

```
rpart.plot(rpart(formula = Churn ~., data = train,  
                 method = "class", parms = list(split = "information")), extra = 100)
```

**Output:**



Showing confusion matrix

**Code:**

```
confusion_matrix <- table(test_pred_info, test$Churn)
```

```
print(confusion_matrix)
```

**Output:**

```
test_pred_info    0     1
0 56701 1294
1   395 73484
```



Showing Precision, recall and F1 score

**Code:**

```
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
precision <- diag(confusion_matrix) / rowSums(confusion_matrix)
recall <- diag(confusion_matrix) / colSums(confusion_matrix)
f1_score <- 2 * (precision * recall) / (precision + recall)

report <- data.frame(Class = rownames(confusion_matrix),
                     Precision = precision,
                     Recall = recall,
                     F1_Score = f1_score)

print(report)
```

**Output:**

```
  Class Precision  Recall  F1_Score
0      0 0.9776877 0.9930818 0.9853247
1      1 0.9946534 0.9826954 0.9886383
```

## Summary

	accuracy	Precision	Recall	F1	TP	TN	FP	FN	No. of rules
Gini	97%	0 95% 1 99%	0 99% 1 96%	0 97% 1 97%	56701	71965	395	2813	6
Information	98%	0 97% 1 99%	0 99% 1 98%	0 98% 1 98%	56701	73484	395	1294	8

**Accuracy:** This indicates the overall correctness of the decision tree's predictions on the dataset. The Gini-based tree has an accuracy of 97%, while the Information Gain-based tree has an accuracy of 98%. This suggests that both trees are performing well.

**Precision:** Precision is a measure of how many of the positive predictions made by the model are actually correct. For class 0 and class 1, both trees have a precision of 99%, meaning that the vast majority of the positive predictions are accurate for both classes.

**Recall:** Recall measures how well the model is able to identify all positive instances. For class 0, both trees have a recall of 95%, while for class 1, the recall is 96% for the Gini-based tree and 98% for the Information Gain-based tree. This indicates that the Information Gain-based tree is better at correctly identifying positive instances of class 1.

**F1-Score:** The F1-Score is a balanced measure that takes both precision and recall into account. For class 0, both trees have an F1-Score of 99%, while for class 1, the F1-Score is 97% for the Gini-based tree and 98% for the Information Gain-based tree. Again, the Information Gain-based tree has a slight advantage in class 1.

**True Positives (TP):** The number of instances that are correctly predicted as positive.

**True Negatives (TN):** The number of instances that are correctly predicted as negative.

**False Positives (FP):** The number of instances that are predicted as positive but are actually negative.

**False Negatives (FN):** The number of instances that are predicted as negative but are actually positive.

**Number of Rules:** This column indicates the number of rules extracted from each decision tree. The Gini-based tree generates 6 rules, while the Information Gain-based tree generates 8 rules. This suggests that the Information Gain-based tree have a more complex structure with more conditions for making decisions. As a result, we can say the Gini based tree is computationally less expensive

In summary, both decision trees seem to perform well, with the Information Gain-based tree having a slightly better accuracy, recall, and F1-Score for class 1 predictions. But the Gini based tree is computationally less expensive compare to information gain-based tree